

# Introduction to Neural Networks

**AE4-350** Bio-inspired intelligence and learning for aerospace applications

Erik-Jan van Kampen

Delft University of Technology

May 1, 2019

# Outline

- 1 Introduction
- 2 Viewing neural networks from a non-mathematical perspective
- 3 Coupling visual perspective to mathematical description
- 4 Off-line training
- 5 (Re-)defining your neural network based on training performance
- 6 On-line training
- 7 Advanced neural network structures and training algorithms

# Learning objectives

After this lecture, the student can:


- Explain the relationship between biological neural networks and artificial neural networks.
- Describe the types of neural networks and their applications.
- Design and construct his/her own artificial neural network for a given task.



 NASA Dryden Flight Research Center Photo Collection  
<http://www.dfrc.nasa.gov/Gallery/PhotoIndex.html>  
NASA Photo: EC03-0251-2 Date: August 27, 2003 Photo By: Jim Ross  
NASA Dryden's highly-modified F-15B aircraft, tail number 837, serves as an Intelligent Flight Control System (IFCS) research testbed aircraft.

# Introduction

- First appearance of the artificial neural network in the late 1800's but the modern era started in 1943 with the work of McCulloch and Pitts.



# Introduction

- Artificial neurons are a mathematical representation of the neurons in the human brain.




Figure: Biological neuron and the mathematical equivalent.

# Introduction

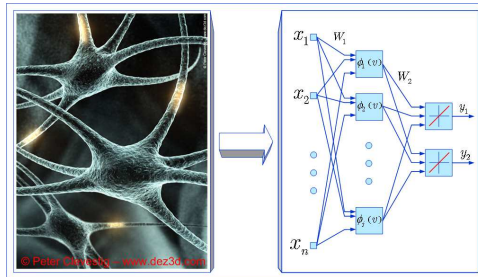


Figure: Mathematical representation of a network of neurons

- Initially the use of neural networks was limited due to difficulties in setting the weights
- Introduction of the error back-propagation learning method (1986, Rumelhart, Hinton, and Williams) created an efficient way in defining weights and gave a new incentive to the field on neural networks.

# Introduction

- In the last decades neural networks have been used in a lot of different applications:
  - Pattern recognition
  - Speech recognition
  - Adaptive control
  - Stock market forecasting
  - Weather forecasting
  - Aerodynamic model identification
  - Games: checkers, backgammon, chess, ...
  - ...
- The method of incorporating neural networks for system identification and control can vary...

# Introduction


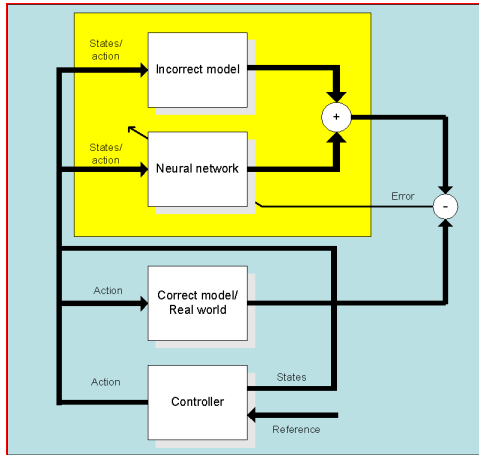


Figure: Neural networks for complete system identification

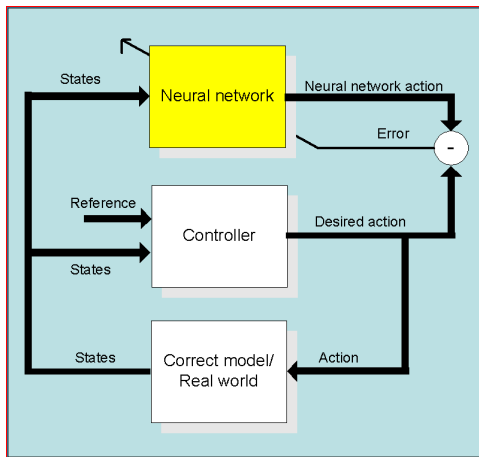


# Introduction



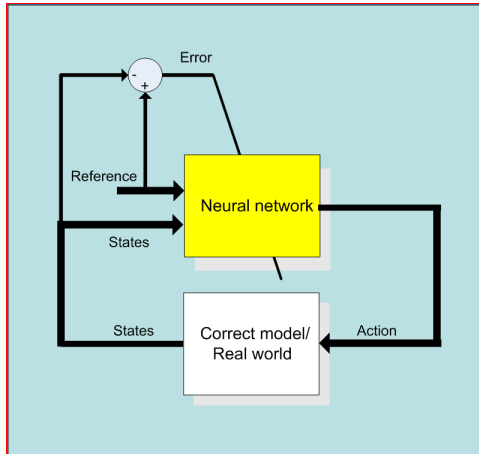
**Figure:** Neural networks for partial system identification. The neural network is used to correct a partial model or a linear model. This is one way of incorporating knowledge in system identification.

# Introduction



**Figure:** Neural networks for control. Training of the network is performed using a supervisor. The supervisor provides the error in the control output. With this error the network weights are updated.

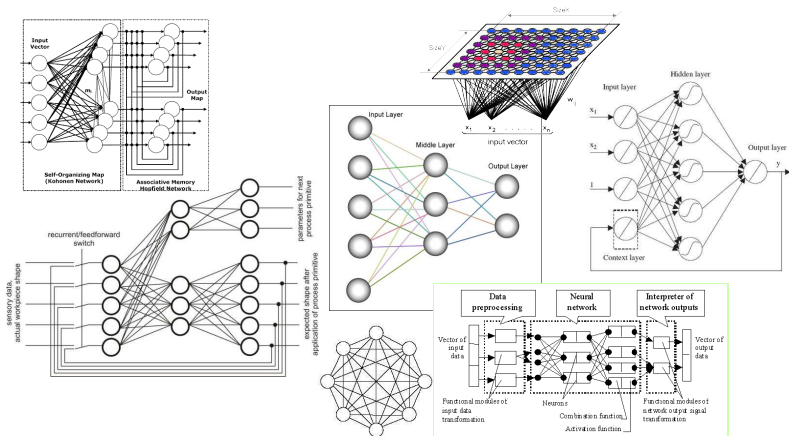
# Introduction



**Figure:** Neural networks for control. Unsupervised learning of the neural network. The network uses the difference between the reference and the state to adapt its weights, i.e. adapt the control law.

# Introduction

- Many different network structures exist today ranging from simple to extremely complex networks.



- Aircraft Detection using Deep Convolutional Neural Network for Small Unmanned Aircraft Systems. Sunyou Hwang, Jaehyun Lee, Heemin Shin, Sungwook Cho and David Hyunchul Shim. AIAA SciTech conference 2018.



# Introduction



# Introduction

- All neural network types are designed for a specific task: there is not a single network type which is optimal for all tasks.
- If you are given a task ...
  - which type of neural network is the most optimal?
  - what activation function must one use?
  - how many neurons are needed?
  - how are the neurons linked between layers?
  - which input and outputs are required?
  - what training algorithm must be used?
  - how can one collect the required training/testing data?
  - ...

# Introduction


Approach to achieve the learning goals:

- Static input-output mappings.
- Neural networks from a visual perspective (solving a puzzle).
- Linking the visual perspective to the mathematical representation and network definition.
- Discussion on training the networks offline.
- Discussion on training the networks online.
- Introduction into more complex neural networks.
- Viewing some applications of neural networks.



# Visual inspection of neural networks

- Neural network design and optimization is much like puzzling.
- This will be demonstrated for a simple neural network



using a MATLAB demo...

# Visual perspective on NN's

The demonstration has shown ...

- that a neural network output is a collection of weighted activation functions.
- the placement and shaping of the neurons can be difficult.
- activation function can be any desired function.
- the performance of the network depends on the chosen activation functions although almost perfect fits can be made with enough neurons.
- for any IO mapping there is an optimal set of activation functions such that the neural network contains the lowest amount of adaptable parameters to obtain a particular cost function value.

# Visual perspective on NN's

## Important questions


- How can one determine the optimal set of activation functions?
- What can one do if no information is available regarding the IO mapping?

Theorem on neural networks:

Cybenko (1989):

*A feedforward neural net with at least one hidden layer with sigmoidal activation functions can approximate any continuous nonlinear function  $\mathbb{R}^p \rightarrow \mathbb{R}^n$  arbitrarily well on a compact set, provided that sufficient number of hidden neurons are available.*


# Visual perspective on NN's




**Figure:** General structure of a feedforward neural network with a single hidden layer. Summation of inputs per neuron is performed inside the neuron before the activation function is applied.

# Visual perspective on NN's


Demonstration of approximation power:




(a) Real IO mapping




(b) NN 25 neurons




(c) NN 50 neurons



(d) NN 75 neurons




(e) NN 100 neurons




(f) NN 200 neurons

# Visual perspective on NN's


Demonstration of approximation power: error plots




(a) NN 25 neurons




(b) NN 50 neurons



(c) NN 75 neurons



(d) NN 100 neurons



(e) NN 200 neurons

# Mathematical description

General mathematical description of a single hidden layer feedforward neural network:

$$\begin{aligned}y_k &= \phi_k(v_k); \\v_k &= \sum_j w_{jk} y_j = \sum_j w_{jk} \phi_j(v_j) \\v_j &= \sum_i w_{ij} x_i = \sum_i w_{ij} \phi_i(v_i)\end{aligned}$$


Simplified mathematical description:

$$\left. \begin{aligned}y_k &= v_k \\v_k &= \sum_j w_{jk} \phi_j(v_j) \\v_j &= \sum_i w_{ij} x_i\end{aligned} \right\} y_k = \sum_j w_{jk} \phi_j \left( \sum_i w_{ij} x_i \right)$$

# Mathematical description

- The input weights  $w_{ij}$  and output weights  $w_{jk}$  determine the position and the shape of the activation function.
- This can be shown by a simple demonstration for a single hidden neuron neural network...


$$y = w_{1k}\phi_j(w_{11}x + w_{01}) + w_{0k}$$




*Worksheet!*




# Mathematical description




(a) FF structure




(b) Adaptation of  $w_{01}$



(c) Adaptation of  $w_{11}$



(d) Adaptation of  $w_{0k}$



(e) Adaptation of  $w_{1k}$


# Mathematical description

- The feedforward neural networks with sigmoidal activation functions can approximate any function.
- It is one of the most popular networks.
- Another very popular type of network is the radial basis function neural network:


$$\left. \begin{aligned} y_k &= v_k \\ v_k &= \sum_j w_{jk} \phi_j(v_j) \\ v_j &= \sum_i w_{ij} (x_i - c_{ij})^2 \end{aligned} \right\} y_k = \sum_j w_{jk} \phi_j \left( \sum_i w_{ij} (x_i - c_{ij})^2 \right)$$

The activation function looks like ...


# Mathematical description




(a)  $W_{ij} = \{10, 5\}$ ,  $C_{ij} = \{0, 0\}$




(b)  $W_{ij} = \{10, 5\}$ ,  $C_{ij} = \{1, 0\}$




(c)  $W_{ij} = \{20, 5\}$ ,  $C_{ij} = \{0, 0\}$



(d) Top view



(e) Top view



(f) Top view

# Mathematical description



## RBF network

- Local character
- Easy to model local complexity
- Bad generalization?
- High computation load
- Easy optimization



## FF network

- Global character
- Easy to model global trends
- Good generalization?
- Low computation load
- Difficult optimization



# Mathematical description

## Summary

- Two commonly used neural networks introduced
  - Neural network output is a summation of activation functions
  - Characteristics of IO mapping should be reflected by the activation functions
  - Shape of the activation functions adjusted by network weights
  - Information stored in the network weights
  - IO mapping complexity together with choice in activation function determines the required number of neurons.
- 
- In the following slides, static IO mappings in combination with standard feedforward networks and RBF networks are considered.
  - Next step: training networks ...



# Training networks

## Start-off point

- Single hidden layer feedforward neural network (fixed structure)

$$y_k = \sum_j w_{jk} \phi_j \left( \sum_i w_{ij} x_i \right)$$

- Static IO mapping, fixed set of data  $(x, d)$ .



# Training networks

Aspects off training neural networks:

- Network structure
- Network weight initialization
- IO data pre-processing
- Adaptation rules
- Stopping conditions

For now consider ...

- Random initialization
- No data pre-processing


Next: error back propagation...

# Error-back propagation

- Network performance measured with a squared cost function:

$$E = \frac{1}{2} \sum_q \sum_k (d_{k,q} - y_{k,q})^2$$

- The error-back propagation approach tries to minimize the cost function by traveling down the slope:





# Error-back propagation


- Updates performed using partial derivatives: step taken in the negative gradient direction.
- Update law for the network weights given by:

$$w_{t+1} = w_t + \Delta w; \Delta w = -\eta \frac{\partial E}{\partial w_t}$$

- Update law depends on ...
  - 1 ... partial derivatives
  - 2 ... learning rate parameter  $\eta$
- Partial derivative can be computed in a structured way using the chain rule: back-propagation of the errors


# Error-back propagation

- First a forward computation step is performed to compute the neuron inputs ( $v$ ) and outputs ( $y$ ) (consider a single IO data point)
- Then the output errors are computed ( $e_k$ )
- Finally the cost function dependencies are propagated from right to left...



# Error-back propagation


- Step 1: dependencies w.r.t. network outputs



$$E = \sum_k \frac{1}{2} (d_k - y_k)^2$$
$$\frac{\partial E}{\partial y_k} = \frac{\partial E}{\partial e_k} \frac{\partial e_k}{\partial y_k} = e_{k,q} - 1$$

# Error-back propagation


- Step 2: dependencies w.r.t. neuron inputs



$$\begin{aligned} y_k &= \phi_k(v_k) \\ \frac{\partial y_k}{\partial v_k} &= \frac{\partial \phi_k}{\partial v_k} \end{aligned} \quad \text{linear activation function} \rightarrow \frac{\partial y_k}{\partial v_k} = 1$$

# Error-back propagation

- End goal 1: dependencies w.r.t. output weights




$$v_k = \sum_j w_{jk} y_j \rightarrow \frac{\partial E}{\partial w_{jk}} = \frac{\partial E}{\partial y_k} \frac{\partial y_k}{\partial v_k} \frac{\partial v_k}{\partial w_{jk}} = e_{k,q} \cdot 1 \cdot \frac{\partial \phi_k}{\partial v_k} \cdot y_j$$

$$\frac{\partial v_k}{\partial w_{jk}} = y_j$$

# Error-back propagation


- Step 3: dependencies w.r.t. neuron output



$$\begin{aligned} v_k &= \sum_j w_{jk} y_j \rightarrow \frac{\partial E}{\partial y_j} = \sum_k \frac{\partial E}{\partial y_k} \frac{\partial y_k}{\partial v_k} \frac{\partial v_k}{\partial y_j} \\ \frac{\partial v_k}{\partial y_j} &= w_{jk} \end{aligned}$$

# Error-back propagation


- Step 4: dependencies w.r.t. neuron input



$$\begin{aligned} y_j &= \phi_j(v_j) \\ \frac{\partial y_j}{\partial v_j} &= \frac{\partial \phi_j}{\partial v_j} \end{aligned}$$

# Error-back propagation

- End goal 2: dependencies w.r.t. input weights



$$\begin{aligned}
 v_j &= \sum_i w_{ij} y_i & \frac{\partial E}{\partial w_{ij}} &= \sum_k \frac{\partial E}{\partial y_k} \frac{\partial y_k}{\partial v_k} \frac{\partial v_k}{\partial y_j} \frac{\partial y_j}{\partial v_j} \frac{\partial v_j}{\partial w_{ij}} \\
 \frac{\partial v_j}{\partial w_{ij}} &= y_i & &= \sum_k e_{k,q} - 1 \cdot \frac{\partial \phi_k}{\partial v_k} \cdot w_{jk} \cdot \frac{\partial \phi_j}{\partial v_j} \cdot y_i
 \end{aligned}$$



# Error-back propagation

- Complete partial derivatives (for all data points):

- Output weights:

$$\frac{\partial E}{\partial w_{jk}} = \sum_q \frac{\partial E}{\partial e_{k,q}} \frac{\partial e_{k,q}}{\partial y_{k,q}} \frac{\partial y_{k,q}}{\partial w_{jk}} = \sum_q e_{k,q} \cdot -1 \cdot y_j$$

- Input weights:

$$\begin{aligned} \frac{\partial E}{\partial w_{ij}} &= \sum_q \left[ \sum_k \left\{ \frac{\partial E}{\partial e_{k,q}} \frac{\partial e_{k,q}}{\partial y_{k,q}} \frac{\partial y_{k,q}}{\partial y_{j,q}} \frac{\partial y_{j,q}}{\partial v_{j,q}} \frac{\partial v_{j,q}}{\partial w_{ij}} \right\} \right] \\ &= \sum_q \sum_k e_{k,q} \cdot -1 \cdot w_{jk} \frac{\partial \phi_{j,q}}{\partial v_{j,q}} x_i \end{aligned}$$


- If multiple layers are present than use the same method.

# Error-back propagation

- Now that the derivatives are known the update of the weights can be performed:

$$w_{t+1} = w_t + \Delta w; \Delta w = -\eta \frac{\partial E}{\partial w_t}$$

- The learning rate is used to determine the step size
- Together with the initialization point they determine the final approximation performance:




# Error-back propagation


## Basic learning algorithm: fixed learning rate

- Initialization
  - LOOP (number of epochs)
    - Compute cost function value  $E_t$  and weight update  $\Delta W$  for current set of weights  $W_t$
    - Perform update of weights and compute new cost function value  $E_{t+1}$
    - If  $E_{t+1} < E_t$  then accept changes else stop loop
  - END LOOP
- 
- Small learning rate  $\rightarrow$  slow convergence to local minimum (derivatives zero)
  - Large learning rate  $\rightarrow$  fast convergence, but end may not be a true minimum (derivatives non-zero)

# Error-back propagation



(a) Approximation



(b) Performance

**Figure:** Influence of the learning rate on the performance time history during training

# Error-back propagation

## Learning algorithm: adaptive learning rate

- Initialization
- LOOP (number of epochs)
  - Compute cost function value  $E_t$  and weight update  $\Delta W$  for current set of weights  $W_t$  and  $\eta_t$
  - Perform update of weights and compute new cost function value  $E_{t+1}$
  - If  $E_{t+1} < E_t$  then accept changes and increase learning rate  $\eta_{t+1} = \eta_t * \alpha$  else do not accept changes and decrease learning rate  $\eta_{t+1} = \eta_t * \alpha^{-1}$ .
  - Stop loop if partial derivatives are (nearly) zero.
- END LOOP

# Error-back propagation

- The previous training algorithms are all first order gradient-descent algorithms (a.k.a. steepest descent)
- There are second order methods available of which the most commonly used is the Levenberg-Marquardt method:

$$\mathbf{w}_{t+1} = \mathbf{w}_t - (\mathbf{J}^T \mathbf{J} - \mu \mathbf{I})^{-1} \mathbf{J} \mathbf{e}; \quad \mathbf{J} = \frac{\partial \mathbf{y}}{\partial \mathbf{w}_t}$$


- If  $\mu$  is larger then the LM method reduces to a first order steepest-descent method:

$$\mathbf{w}_{t+1} \approx \mathbf{w}_t - (\mu \mathbf{I})^{-1} \mathbf{J} \mathbf{e} = \mathbf{w}_t - \mu^{-1} \mathbf{J} \mathbf{e}$$


- If  $\mu$  is small then the LM method reduces to a pure Newton step:

$$\mathbf{w}_{t+1} \approx \mathbf{w}_t - (\mathbf{J}^T \mathbf{J})^{-1} \mathbf{J} \mathbf{e}$$

# Error-back propagation



(a) Approximation



(b) Performance

**Figure:** Influence of the training algorithm on the performance time history during training

Levenberg-Marquardt is much more efficient and shows a faster convergence.

# Error-back propagation

## 1 order method

- Low computation cost
- Numerically stable
- Slower convergence
- Gentle weight updates
- Easy implementation



## 2 order method


- Higher computation cost
- Possibly unstable
- Fast convergence
- Aggressive weight updates
- More difficult implementation




# Initialization

- As mentioned previously, the initialization of the weights has a large influence
- For feedforward neural networks with sigmoidal activation function random initialization is used: small weight values means that the first order derivative is non-zero in the entire input space.
- For the RBF neural networks one can do the same but also uniform distribution can be performed (see next slide).
- Optimal initialization requires knowledge of the true IO mapping.
- Common practice is to perform the training several time, each time with different initial weights.


# Initialization




(a) Uniform



(b) Random



(c) IO mapping



(d) IO mapping

# Neural network training

## Summary

- Several basic training algorithms explained
  - Error-back propagation algorithms most commonly used
  - Fast convergence, aggressive steps vs slower convergence, gently steps
  - Initialization has a large influence on the final performance
- 
- Next step: defining your neural network in a structured way ...

# Defining your network


- Start-off point:
  - Input and output are set
  - Activation functions are set
  - Static IO mapping
  - IO data available

## Question


How many neurons do you need for optimal performance?

- Approach: Iteration based on performance
  - Perform training phase for a given number of neurons and several weight initializations.
  - Increase the number of neurons and if the performance increase is large enough: add more neurons.
- Pitfalls: Over-fitting, loss of generalization
- An example ...

# Defining your network



(a) Performance




(b) Output for 30 hidden neurons

Figure: Example of the over-fitting problem

# Defining your network


- Common approach to prevent over-fitting: separation of training and test data
- Training algorithm stops if the cost function value, using both data sets, is non-decreasing!




- Common ratios between training and test data: 90/10, 80/20.
- The test data will pose an upper limit on the number of neurons!
- An example ...

# Defining your network

- 90% training data, 10% test data, per network 10 initializations:



(a) Performance



(b) Output for two networks

Figure: Example of using test data to prevent over-fitting

# Defining your network

- For RBF networks there are more possibilities with optimization since the effects per neuron are local
- A lot of pruning and allocation methods exist
- Principle of training/test data still applies
- Routines are available in the MATLAB neural network toolbox



# Mathematical description

## Summary up to now


- Covered all basic aspects of off-line training for static IO mappings:
  - Activation function selection
  - Optimal number of neurons
  - Weight initialization
  - Training algorithms
- Two most commonly used networks treated: FFNN and RBFNN
- Creating the optimal network requires knowledge of the IO mapping
- No universal way of finding the optimal network

What remains to be treated:

- On-line training
- Discussion on more complex neural networks
- Interval analysis
- Assignment

# On-line training

- Required when the IO mapping is time-varying:
  - Aerodynamic model identification in case of failures
  - UAV motor dynamics
  - Reinforcement learning
- On-line training usually required in a closed loop system identification setting



# On-line training

- Problems with on-line training:
  - Gathering enough IO data
  - Loss of information during training
  - Problems of network optimization (activation functions, number of neurons, structure)
  - Defining inputs
  - Convergence issues
- Approach taken in the following slides
  - Random initialized neural network (mimic previous IO mapping)
  - Sweep through input space to gather IO data
  - Sequential training

# On-line training

- Example of the problems with on-line training ... (movie)
- Problems with FFNN: global behavior of activation functions
- Problem can be reduced by using:
  - Time-window with batch updating
  - Gentle updates: steepest descent instead of Levenberg-Marquardt
  - Anti-recency points ... (movie)
  - Store previous IO data points
  - Activation function with a local character like RBF ... (movie)

# On-line training

## Summary

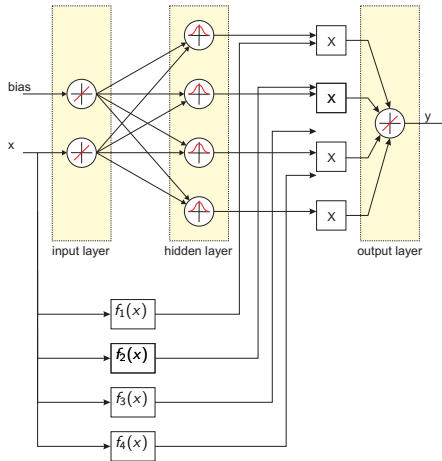
- When dealing with off-line identification neural networks are very suitable
- Optimization of neural networks is a non-linear optimization problem making neural networks less suitable for on-line system identification.
- Neural networks with 'local' activation functions such as RBF perform better in online learning.
- There will always be a trade-off between remembering previously learned information and learning new information.

# Advanced neural network structures

- Theory shows that all IO mappings can be approximated to any desired accuracy with a single hidden layer feedforward neural network with sigmoidal activation functions.
- Practice has shown that the required number of neurons can be huge.
- Other activation functions are better suited or other network definitions are better suited.
- One very powerful group of networks uses the combination of local activation functions and polynomials...

# Advanced neural network structures


- Example of network layout:




- Multiple layouts possible: most variations in the definition of the RBF neurons.

# Advanced neural network structures

- The RBF neurons are used to defined sub-domains in the input-space.
- The polynomials approximate the IO mapping on each sub-domain.




(a) RBF neurons




(b) Polynomials



# Advanced neural network structures



(a) Output per neuron



(b) Total network output

- That the IO mapping is continuous up to any degree (gradient descent methods applicable).
- RBF neurons 'blend' the polynomials between sub-domains.
- Local complexity can be adjusted by setting polynomial degree per neuron.

# Advanced neural network structures

- Optimization of RBF neurons can be made based on allocation and pruning techniques.
- Optimization of polynomial order can be made based on performance on the sub-domains.
- Network comparable to fuzzy neural networks or multivariate splines (lectures of dr.ir. Coen de Visser)
- Note that the optimization of the polynomial coefficients is a convex optimization problem once the RBF weights are fixed.

# Advanced neural network structures

- Other variation possible: Use of more elaborate RBF neuron definitions

$$u_t = \frac{1}{\sum_j \left( \frac{\|\mathbf{x} - \mathbf{v}_t\|}{\|\mathbf{x} - \mathbf{v}_j\|} \right)^{\frac{2}{m-1}}}$$

- More complex sub-domain definitions




Figure 5: Output per neuron (or Degree of Membership functions)

# Advanced training algorithms

- Neural network training (partially) non-linear optimization problem.
- Outcome of gradient-based algorithms dependent on weight initialization.
- More complex non-linear optimization algorithms are available to 'solve' the problem
  - Line search methods
  - Advanced gradient methods (better convergence, higher probability of optimal solution)
  - Genetic algorithms
  - etc...
- Previous methods can never give the guarantee that the global optimum or *all* global optima are found!
- Methods that do have this guarantee
  - Lipschitz bound based algorithms
  - Interval analysis based algorithms

# Important notes on neural networks

- Each task has a unique network definition which is optimal (least amount of free parameters).
- Most complex networks can be broken down into the elements described in this lecture.
- Many learning algorithms available: start with gradient descent methods.
- Neural networks less suited for on-line applications, but if applied be aware of the recency effect (usually less neurons is better).
- When working with neural networks: if knowledge is available then use it to define your network and try to keep the network as small as possible! Less neurons means faster learning and fewer local minima!

# Introduction to Neural Networks

**AE4-350** Bio-inspired intelligence and learning for aerospace applications

Erik-Jan van Kampen

Delft University of Technology

May 1, 2019